

GAE Feedback


Distributed Systems 2016-2017

Key Attention Points

- 1. Persistence: GAE Datastore using JPA
 - Structure of entities in Entity Groups
 - Use of JPQL
 - Transactions / transactional behaviour
- 2. Indirect Communication
 - Adapt from Direct Communication
 - Use of Task Queues, reason about data flow
- 3. Potential State Inconsistencies (theoretical exercise)
 - Why? How to avoid?
 - How to minimize impact on performance?

1. Persistence

- **Entity Group (EG):** entity relations as hierarchical **tree-structure**
 - one-to-many == parent-child relationship
 - every entity has at max one parent entity (tree)
- **Easiest solution in car rental agency**
 - All entities in a single Entity Group
 - No unmanaged relations between those entities
 - E.g. Manual lookup by UID (Primary Key)




1. Persistence (cont.)

- Google App Engine: JPQL
 - No JOINS supported
 - Use only to query on one entity
- Transactions
 - Each EntityManager session == atomic commit
 - With or without JTA
 - Modification within of a single Entity Group
 - 2 Options for transactional confirming quotes
 - At application level
 - Process all quotes for one CRC within one transaction at a time
 - Undo successful reservations if not all quotes could be confirmed
 - At the database level (requires cross-group transactions (XGT))
 - Process all quotes at once
 - Failure to confirm all quotes will trigger automatically rollback

1. Persistence: User Profiles

- User Profiles
 - Independent from car reservations
 - No relation to other profiles
- Entity Group model as depicted in figure
 - Simple modifications will suffice
 - Any change on a single entity group guarantees ACID
 - No need for cross-group transactions

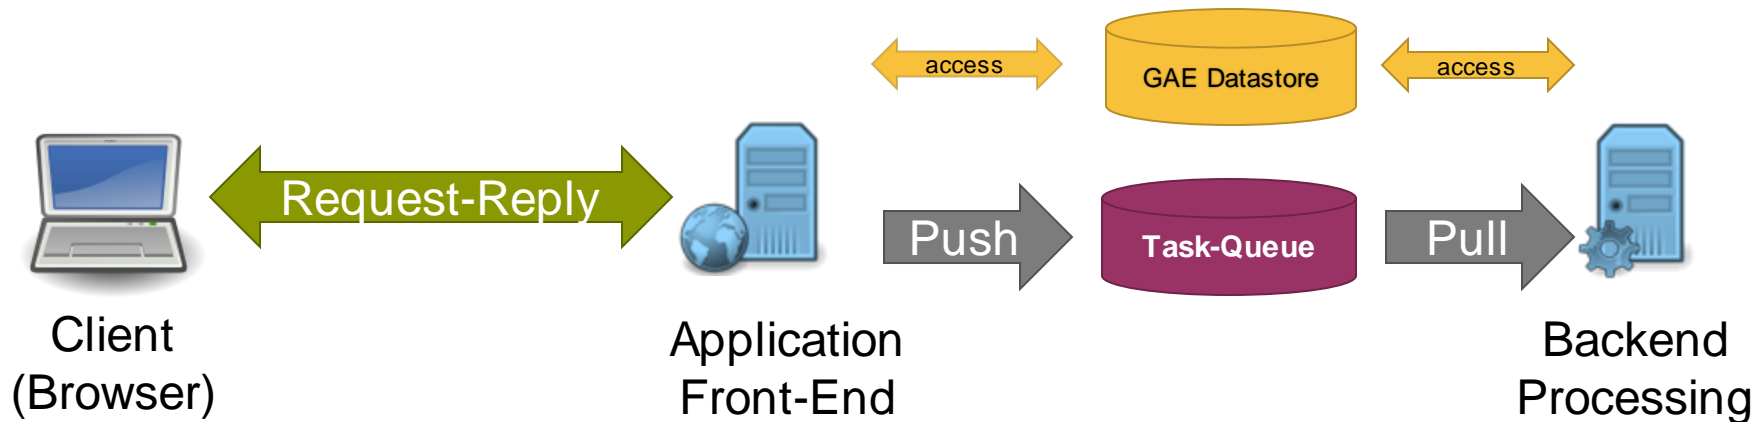


2. Indirect Communication

■ The 3 roles involved

■ Data flow:

- Client (Browser) \leftarrow direct \rightarrow Font-end Service
- Font-end Service – via Queue \rightarrow Back-end Service (Worker)



KEY

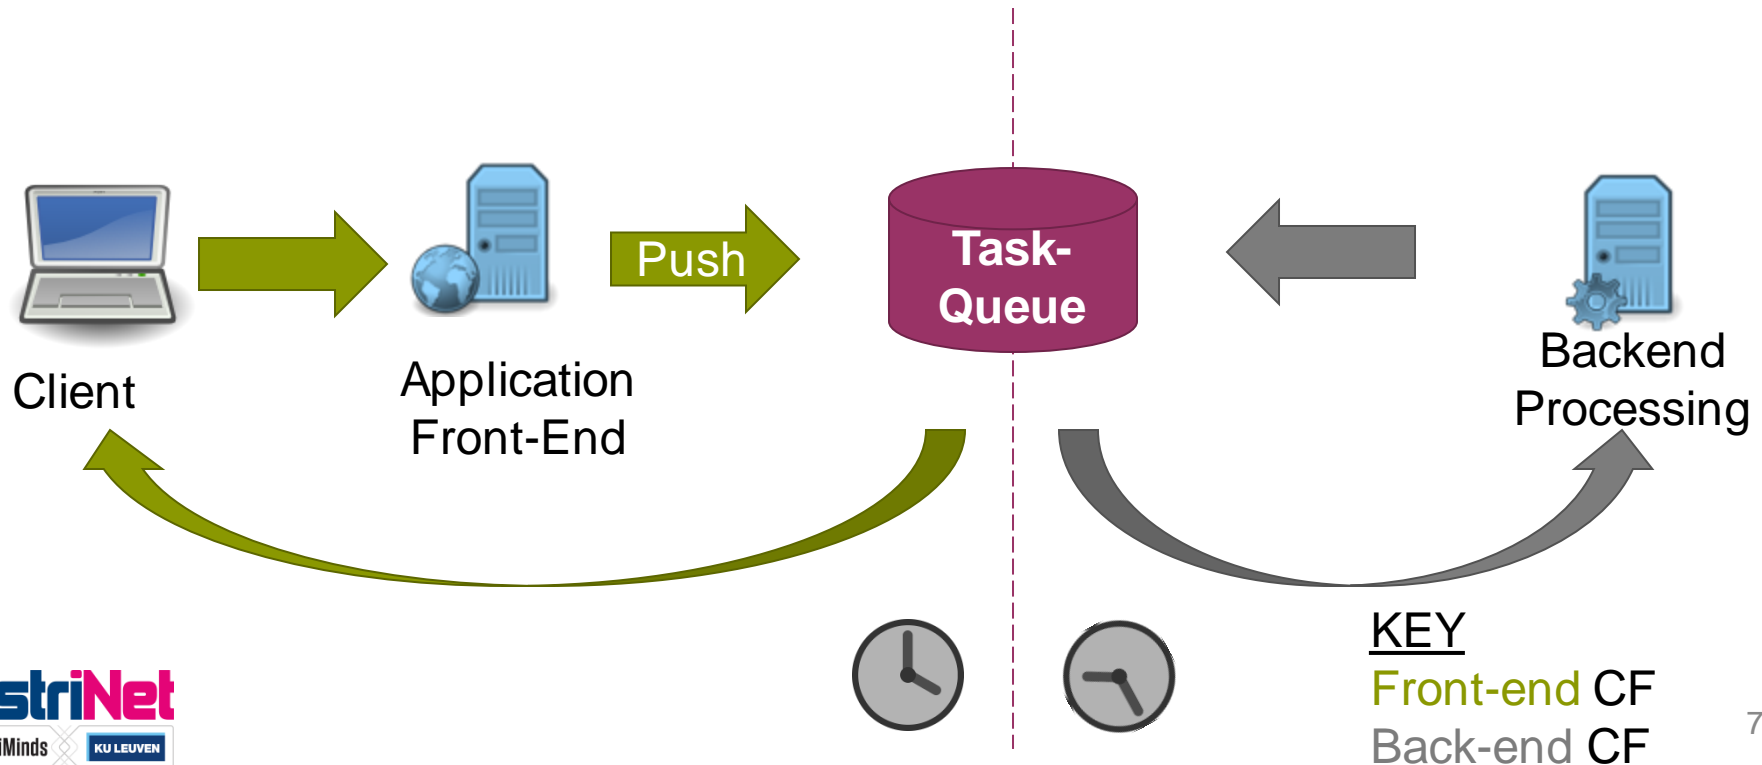
Direct Communication

Indirect Communication

2. Indirect Communication

Control flow (CF):

- Front-end service triggered by Client
- Back-end service triggered by GAE infrastructure
- (1) and (2) are **independent in time**



2. Indirect Communication

- **Feedback Channel: Do not**
 - Send msgs directly from back-end to front-end
 - client may be gone at time of message
 - wait with front-end until back-end completes
- **Examples**
 - no shared memory / cache (shared „global“ variables)
 - no Channel API
- **Backend workers use separate machines**
 - in own process (\Leftrightarrow own thread)
 - no shared memory
 - shared „global“ static variables > won't work

2. Indirect Communication

- **Feedback channel** from Worker to Client
 - report about success AND failure
 - otherwise indistinguishable : not-yet-processed or failed
- Client may have **multiple** unconfirmed **quotes waiting**, or may have sent a rental order twice.
 - feedback should uniquely refer to quotes (e.g. using **orderID**)

3. Potential State Inconsistencies

- **State changes** (creation of reservations) happen only at the Worker role and not at the front-end.
 - Keyword **,synchronized'** for quote-enqueue function **no effect** on potential state inconsistencies.
- Multiple workers process tasks queue **in parallel** (default)
 - This is where potential inconsistencies are rooted
- Each worker is a separate **process** (separate JVM)
 - thread-based monitors (i.e. **,synchronized'**) **cannot help**

3. Potential State Inconsistencies

- Bogus behavior (double booking) possible?
- With Cross-Group Transactions (XGT)
 - No, locking at the database-level
- Without XGT
 - Depends on GAE implementation of optimistic concurrency control
 - At persistence layer: no, concurrency control sufficient
 - At middleware layer: yes, different processes not aware of each other >> solution addresses this case

3. Potential State Inconsistencies

■ One solution

- Set #worker per queue = 1
- To increase parallelism
 - Maintain one queue per company
(**Note**: only if client will book from single CRC)

■ NOT

- Create a queue per CarType or per Quote
 - → loss of all-or-nothing semantics
- Create a queue for every task
 - japodizes state consistency even more (parallelism $\rightarrow \infty$)

Questions?